

Osaid Khan

+1 (945) 304-5781 | khanosaid726@gmail.com | www.linkedin.com/in/-osaid-khan/ | github.com/ozzyozbourne

EDUCATION

Pace University <i>Master of Science (MS) in Computer Science (GPA - 4.0/4.0)</i>	New York City, NY <i>Jan 2024 - Dec 2025</i>
Dr. A.P.J. Abdul Kalam Technical University <i>Bachelor of Technology (B.Tech) in Computer Science</i>	Noida, India <i>Jul 2016 - Sep 2020</i>

TECHNICAL SKILLS

Languages: Java, JavaScript/TypeScript, Rust, Python
Databases: PostgreSQL, MongoDB, Redis
Others: Git, CI/CD, Node.js, Google Cloud, Kafka
AWS Services: EC2, RDS, S3, Lambda, CloudWatch, Step Functions, CloudFormation, CloudFront, API Gateway
Frameworks: FastAPI, FastMCP, LangGraph, LangChain, LangSmith, PyTorch, React, Sprint Boot
AI: Claude code, Codex, Pi coding-agent

PROFESSIONAL EXPERIENCE

Software Engineer Sperse	Phoenix, AZ <i>Feb 2026 - present</i>
------------------------------------	--

- Developed a cloud-native interaction layer that translates natural-language intents into **safe, validated tool calls** across **1000+ platform endpoints**, exposing **RESTful APIs** aligned with **.NET (8/9/10)**, **C#**, and **ASP.NET Core**
- Engineered **background/worker services** (queues, retries, idempotency) and reusable **data access** modules using **EF Core** patterns for **PostgreSQL** and **MS SQL Server** with **Data Annotations** and **FluentValidation** guardrails
- Optimized routing and multi-agent concurrency with API-aware templates and SDK generation using **AI-assisted development tools** (**GitHub Copilot** and **Claude Code**), then operationalized releases through **CI/CD pipelines**, **DevOps practices**, **Agile methodologies**, and strict **pull request workflows** with **xUnit**, **NUnit**, and **MSTest** test gates; packaged services for **Azure services**, **Docker**, and **Kubernetes (AKS)**, reducing latency by **60%** and LLM API costs by **50%**

AI Researcher Pace Artificial Intelligence Lab	New York City, NY <i>Sep 2025 - Dec 2025</i>
--	---

- Conducted research on **StyleGAN** and **Stable Diffusion** to generate industry-standard fashion designs using **CUDA**, **PyTorch**, and **LORA (Low-rank Adaptation)** on Pace University's **HPC cluster** for model training and fine-tuning
- Utilized **vLLM** and **Hugging Face** deployment in scalable cloud environments, validating model behavior through repeatable integration tests and production inference checks
- Conducted workshops on GPU programming with **CUDA**, **Mojo**, and building **AI Agents** with **LangChain**, **LangGraph**, and **LangSmith** for **+50** students, emphasizing collaborative code reviews and technical communication

Software Engineer Qualitest	Noida, India <i>Feb 2021 - Dec 2023</i>
---------------------------------------	--

- Built an internal support chatbot using **LangChain** and **RAG** (Retrieval-Augmented Generation) with a **PostgreSQL** vector store and **RESTful APIs**, reducing average ticket resolution time by **20%** through validated query-to-answer pipelines
- Designed a cloud-agnostic Storage Service using **Python**, **PostgreSQL**, and **Redis** for uploads/downloads across **Google Cloud Storage**, **S3**, and **MinIO** using expiring signed URLs and strict request validation
- Developed **AWS Step Functions** workflows to run long-running imports as **background/worker services**, prevent **AWS Lambda** timeouts, and improve release stability through repeatable CI integration

PROJECTS

Dark Factory - Orchestrator for AI coding CLIs <i>Crossterm, Ratatui, Y-crdt</i>	github.com/ozzyozbourne/dark-factory
---	---

- Enables full ownership of CLI IO via **PTY**, checking and logging each command before execution by bash shell and reinforcing safe execution patterns used in API-centric automation
- Provides **event-driven** agentic loops and distributed **ralph** coordination with **Y-crdt**, broadcasting **real-time** updates that eliminate polling overhead and support reliable deployment automation