

Osaid Khan

+1 (945) 304-5781 | khanosaid726@gmail.com | www.linkedin.com/in/-osaid-khan/ | github.com/ozzyozbourne

EDUCATION

Pace University <i>Master of Science (MS) in Computer Science (GPA - 4.0/4.0)</i>	New York City, NY <i>Jan 2024 - Dec 2025</i>
Dr. A.P.J. Abdul Kalam Technical University <i>Bachelor of Technology (B.Tech) in Computer Science</i>	Noida, India <i>Jul 2016 - Sep 2020</i>

TECHNICAL SKILLS

Languages: Python, Go, Java, JavaScript/TypeScript
AI/Agents: LangGraph, LangChain, FastMCP, LlamaIndex, Claude Code, Codex, LangSmith, Langfuse
Backend & Frameworks: FastAPI, Spring Boot, React
Databases: PostgreSQL, Redis, MongoDB
AWS Services: Lambda, Step Functions, S3, API Gateway, CloudWatch, EC2, RDS, CloudFormation, CloudFront
Others: Docker, Kubernetes, CI/CD, Git

PROFESSIONAL EXPERIENCE

Software Engineer Sperse	Phoenix, AZ <i>Feb 2026 - Present</i>
- Designed a Python and FastAPI multi-agent CRM with LangGraph, LangChain, and Langfuse; main orchestrator handled subagent orchestration across 40+ agents for long-horizon task execution - Refactored 500+ backend operations into typed dispatchers, replacing 60+ per-agent REST tools with structured interfaces for reliable tool use, correctness, and tenant-aware production behavior - Built dynamic tool routing and context management exposing only relevant tool groups, reducing context bloat and improving tool selection accuracy by 90% across ambiguous agent requests - Implemented lazy-loaded sub-agent skills with JWT/tenant/role propagation through every agent/backend call, lowering startup overhead while preserving access control for multi-step planning - Built trace-driven evals and safety gates (jailbreak regex, Lakera Guard, SQLite) replaying production conversations across prompts/model versions to catch unsafe agent behavior before release	

AI Researcher Pace Artificial Intelligence Lab	New York City, NY <i>Sep 2025 - Dec 2025</i>
- Built an agent-evaluation pipeline ensembling LLM judges and reward models to rank multi-step tool use runs, exposing failure modes and reliability gaps before prompt changes ship - Implemented LangGraph eval scoring and trace selection, tuning prompts to increase tool-call success by 50% while improving repeatable model/product feedback - Analyzed model behavior across prompts/traces to connect model capabilities with reliable AI systems, emphasizing correctness, regression detection, and technical communication	

Software Engineer Qualitest	Noida, India <i>Feb 2021 - Dec 2023</i>
- Built end-to-end CI/CD with GitHub Actions, reducing manual deployments and improving developer velocity for production engineering teams - Optimized PostgreSQL with window functions, materialized views, and indexes to minimize latency and support low latency performance at scale - Orchestrated long-running jobs with AWS Step Functions and AWS Lambda, improving recovery across 100K+ weekly executions in high reliability distributed systems	

PROJECTS

Dark Factory - orchestrator for PI agent <i>Elixir, TypeScript</i>	github.com/ozzyozbourne/dark-factory
- Built a distributed coding-agent orchestrator decomposing large requests into dependency-aware DAGs for long-horizon task execution, retries, timeouts, and multi-machine AI agents - Modeled AI-native IDE developer experience, editor intelligence, agent-in-the-loop workflows, real-time code understanding, and sandboxed code execution environments	