

Osaid Khan

[+1 \(945\) 304-5781](tel:+19453045781) | khanosaid726@gmail.com | www.linkedin.com/in/-osaid-khan/ | github.com/ozzyozbourne

Dear Hiring Team,

It is with great pleasure that I submit my candidacy for the **Full Stack Engineer** role with **David AI**. I bring over three years of professional experience building full-stack systems, LLM-powered data pipelines, and scalable production platforms, along with a strong academic foundation as a Master of Science in Computer Science graduate from Pace University, NY.

Through my experience across two engineering roles, I have built and delivered production-grade full-stack features and data processing pipelines that directly reflect the full-stack engineering depth and rapid iteration this position demands. As a Software Engineer at Sperse, I built a TypeScript and React Flow based full-stack visual workflow platform that ships new features to thousands of users daily, using Temporal for durable, resumable execution across long-running agentic pipelines. Alongside it, I shipped LLM-powered full-stack features serving 10,000+ users via a multi-agent CRM platform in Python, LangGraph, LangChain, and FastAPI, routing requests across 40+ specialized sub-agents with tenant-aware access control. To power data collection and real-time discovery, I created a FastMCP-powered server exposing 160+ LLM-callable backend actions as full-stack data collection interfaces with real-time OpenTelemetry monitoring, and built LlamaIndex-powered RAG agents using OpenAI embeddings and Pgvector to derive actionable insights from large-scale documentation. I also wired these systems into GitHub Actions CI/CD pipelines with OpenTelemetry trace-based evals to build, deploy, and evaluate LLM-based features before every production release. Prior to Sperse, at Qualitest, I orchestrated scalable data processing pipelines with AWS Step Functions across 100K+ weekly executions, designed a cloud-agnostic storage pipeline in Java and Spring Boot with PostgreSQL and Redis spanning Google Cloud Storage, S3, and MinIO — and optimized PostgreSQL with window functions and materialized views to cut query latency, surfacing results through interactive React visualizations for business stakeholders. Together, these roles have prepared me to contribute meaningfully to David AI's core full-stack engineering responsibilities and ship features that help researchers and customers derive value from audio data at scale.

My project work has reinforced the systems-thinking and distributed-execution skills central to building scalable real-time data pipelines. I built Dark Factory — a distributed full-stack coding-agent orchestrator in TypeScript that decomposes large requests into dependency-aware task DAGs, executing them across multi-machine AI agents with fault-tolerant scheduling, configurable retries, and state recovery. The engineering challenges I tackled there — parallel execution, reliability under failure, and real-time pipeline observability — map directly to the kind of infrastructure David AI builds to process terabytes of audio data daily. I am deeply motivated to apply these full-stack and LLM deployment skills to the audio AI domain, rapidly learning digital signal processing and audio data characteristics alongside David AI's research team to iterate on the tools that make high-quality audio datasets possible.

I have always been passionate about bringing AI into the real world and building the data layer that makes that possible, and I envision myself contributing my talent to the furthering of audio AI as the gateway to accessible, human-centered intelligence. I would love to discuss my interest in this position with your team and can be reached at +1 (945) 304-5781 or khanosaid726@gmail.com. Thank you for your time and consideration, I look forward to connecting!

Sincerely,
Osaid Khan