

Osaid Khan

[+1 \(945\) 304-5781](tel:+19453045781) | khanosaid726@gmail.com | www.linkedin.com/in/-osaid-khan/ | github.com/ozzyozbourne

Dear Hiring Team,

It is with great pleasure that I submit my candidacy for the **GenAI and Python Developer** position with **EXL**. I bring over three years of professional experience in developing Python-based AI applications, fine-tuning Generative AI models, and building RAG systems, along with a strong academic foundation as a Master of Science in Computer Science graduate from Pace University, NY.

Through my experience across engineering roles, I have built and delivered production-grade AI systems that directly reflect **the specialized expertise in GenAI model optimization and Python-based AI workflow development this role requires**. As a Software Engineer at Sperse, I architected a Python-based GenAI interaction layer that translates natural-language intents into validated tool calls across 1,000+ platform endpoints — requiring rigorous prompt engineering and safe execution across a distributed microservices surface. To keep these systems production-ready, I implemented a full observability stack using Langfuse and LangSmith for tracing and performance tuning, reducing end-to-end latency by 60% and LLM API costs by 50% through targeted concurrency and routing optimizations in an Agile environment. Prior to Sperse, as a Software Engineer at Qualitest, I built an internal support chatbot using Python, LangChain, and RAG with a PostgreSQL vector store, enabling natural language queries over documentation and reducing ticket resolution time by 20%. I also designed a cloud-agnostic Storage Service to automate file operations across AWS and GCP via secure signed URLs, ensuring high availability and system stability. Together, these roles have prepared me to **drive innovation in LLM fine-tuning and the deployment of production-grade AI applications** and drive impact within an agile engineering team.

My academic research experience has given me a rigorous foundation in model development, GPU-accelerated computing, and technical communication — skills that translate directly into **the culture of collaboration and excellence that EXL maintains to help clients lead through digital transformation**. As an AI Researcher at the Pace Artificial Intelligence Lab, I fine-tuned Generative AI models including StyleGAN and Stable Diffusion using PyTorch and LoRA on an HPC cluster. I was responsible for managing the full model development lifecycle — from data pipelines and training runs to inference optimization using vLLM and deployment via Hugging Face — ensuring both performance and secure, efficient model serving. Beyond building the systems myself, I led workshops on building AI Agents with LangChain and RAG for an audience of 50+ students. Designing and delivering those sessions required me to break down complex technical concepts clearly and accessibly — directly mirroring **the technical leadership and support for GenAI use cases that define this position**. This experience has shaped me into an engineer who builds with rigor and communicates with clarity, both of which I would bring to **EXL and its mission to deliver transformative value**.

I have always been passionate about **innovation and collaboration** and envision myself contributing my talent to the furthering of **helping clients transform and lead their industries through cutting-edge AI and analytics**. I would love to discuss my interest in this position with your team and can be reached at +1 (945) 304-5781 or khanosaid726@gmail.com. Thank you for your time and consideration, I look forward to connecting!

Sincerely,
Osaid Khan