

Osaid Khan

+1 (945) 304-5781 | khanosaid726@gmail.com | www.linkedin.com/in/-osaid-khan/ | github.com/ozzyozbourne

Dear Hiring Team,

It is with great pleasure that I submit my candidacy for the **Software Engineer, LLM Infrastructure** role at **Fireworks AI**. I bring over three years of professional experience building scalable backend infrastructure, LLM CI/CD pipelines, and production ML systems, along with a strong academic foundation as a Master of Science in Computer Science graduate from Pace University, NY — and I am genuinely excited by Fireworks' work on the fastest, most scalable inference in the industry.

Through my experience across two engineering roles, I have built and operated production-grade infrastructure that directly reflects **the ML-infrastructure depth this position demands**. As a Software Engineer at Sperse, I built LLM CI/CD pipelines with GitHub Actions that run OpenTelemetry trace-based evals as frameworks and safeguards to ensure best-in-class model quality — flagging latency, accuracy, and safety regressions during code reviews and continuous deployment. I built scalable backend infrastructure for a multi-agent CRM platform in Python and Go with FastAPI on AWS, serving 10,000+ users by routing inference requests across 40+ specialized sub-agents, and I created a FastMCP control-plane server exposing 160+ backend actions to LLMs that powers model-serving across multiple tenants with built-in monitoring and reliability safeguards. I also built LlamaIndex-powered RAG data pipelines using OpenAI embeddings in pgvector, containerized with Docker and Kubernetes for scalable, cost-efficient inference. Prior to Sperse, as a Software Engineer at Qualitest, I designed a cloud-agnostic Storage Service in Java and Spring Boot handling distributed file operations across Google Cloud Storage, S3, and MinIO, orchestrated long-running distributed jobs with AWS Step Functions to improve reliability and error recovery across 100K+ weekly executions, and optimized PostgreSQL to reduce query latency and improve cost efficiency across compute and storage layers. Together, these roles have prepared me to **contribute to the design of scalable backend services, model serving systems, and reliability improvements across compute, storage, and networking layers**.

I have always been passionate about **low-latency inference and reliable, user-friendly generative AI platforms**, and I envision myself contributing my talent to the furthering of **Fireworks' mission to be the most reliable and user friendly generative AI platform in the world**. I would love to discuss my interest in this position with your team and can be reached at +1 (945) 304-5781 or khanosaid726@gmail.com. Thank you for your time and consideration, I look forward to connecting!

Sincerely,
Osaid Khan