

Osaid Khan

+1 (945) 304-5781 | khanosaid726@gmail.com | www.linkedin.com/in/-osaid-khan/ | github.com/ozzyozbourne

EDUCATION

Pace University <i>Master of Science (MS) in Computer Science (GPA - 4.0/4.0)</i>	New York City, NY <i>Jan 2024 - Dec 2025</i>
Dr. A.P.J. Abdul Kalam Technical University <i>Bachelor of Technology (B.Tech) in Computer Science</i>	Noida, India <i>Jul 2016 - Sep 2020</i>

TECHNICAL SKILLS

Cloud & Platforms: AWS (EC2, RDS, S3, Lambda, CloudWatch, Step Functions, CloudFormation, API Gateway), Google Cloud
Languages: Python, SQL, JavaScript/TypeScript, Rust, Java
DevOps & IaC: CI/CD, Git, CloudFormation, Kafka
Databases: PostgreSQL, MongoDB, Redis
AI & ML: PyTorch, LangChain, LangGraph, LangSmith, Langfuse, FastAPI, FastMCP
Frameworks: React, Spring Boot, Node.js

PROFESSIONAL EXPERIENCE

Software Engineer Sperse	Phoenix, AZ <i>Feb 2026 - present</i>
- Architected platform operations layer enabling self-service execution of 1000+ API endpoints through automated workflows and infrastructure-as-code, reducing time-to-value and improving reliability via governance guardrails - Built observability infrastructure using Langfuse and LangSmith for monitoring, logging, alerting, and tracing, enabling proactive detection, root-cause analysis, and incident management to reduce MTTR - Optimized platform performance and cost transparency through concurrency tuning, reducing latency by 60% and API costs by 50% (FinOps-aligned optimization) while maintaining service levels - Delivered reusable templates and SDK patterns to accelerate API integration, standardizing implementations and improving developer experience and deployment frequency	

AI Researcher Pace Artificial Intelligence Lab	New York City, NY <i>Sep 2025 - Dec 2025</i>
- Operated ML platform infrastructure on HPC cluster using CUDA and PyTorch for model training, fine-tuning, and deployment of StyleGAN and Stable Diffusion models - Implemented MLOps patterns using vLLM for inference and Hugging Face for secure serving, establishing operational standards for model lifecycle management - Enabled platform adoption through self-service tooling and workshops on GPU programming and AI Agents for 50+ students, improving onboarding	

Software Engineer Qualitest	Noida, India <i>Feb 2021 - Dec 2023</i>
- Designed cloud-agnostic Storage Service using Python, PostgreSQL, and Redis across Google Cloud Storage, S3, and MinIO, implementing infrastructure-as-code for environment provisioning and standardization - Built support chatbot using LangChain and RAG with PostgreSQL vector store, reducing ticket resolution time by 20% through automation and enabling self-service - Contributed to Identity & Access Management platform using React, Spring Boot, and AWS, implementing identity/access controls, secrets management, and governance for 30+ applications - Developed automation workflows using AWS Step Functions and Lambda with infrastructure-as-code (CloudFormation) and CI/CD, improving reliability via operational runbooks	

PROJECTS

Dark Factory - Orchestrator for AI coding CLIs <i>Crossterm, Ratatui, Y-crdt</i>	github.com/ozzyozbourne/dark-factory
- Enables platform operations via PTY with command validation and logging before execution, providing observability and governance for distributed CLI sessions - Implements event-driven coordination using Y-crdt for real-time updates across machines, eliminating polling overhead and improving operational efficiency	