

Osaid Khan

+1 (945) 304-5781 | khanosaid726@gmail.com | www.linkedin.com/in/-osaid-khan/ | github.com/ozzyozbourne

EDUCATION

Pace University <i>Master of Science (MS) in Computer Science (GPA - 4.0/4.0)</i>	New York City, NY Jan 2024 - Dec 2025
Dr. A.P.J. Abdul Kalam Technical University <i>Bachelor of Technology (B.Tech) in Computer Science</i>	Noida, India Jul 2016 - Sep 2020

TECHNICAL SKILLS

Frameworks: React, LangGraph, LangChain, LangSmith, FastAPI, PyTorch, FastMCP, Sprint Boot
Languages: JavaScript/TypeScript, Python, Rust, Java
AI: Claude code, Codex, Pi coding-agent
AWS Services: EC2, RDS, S3, Lambda, CloudWatch, Step Functions, CloudFormation, CloudFront, API Gateway
Databases: PostgreSQL, MongoDB, Redis
Others: Node.js, Git, CI/CD, Google Cloud, Kafka

PROFESSIONAL EXPERIENCE

Software Engineer Sperse	Phoenix, AZ Feb 2026 - present
- Lead development of an AI-native interaction layer using TypeScript and Next.js that translates natural-language intents into validated tool calls across 1000+ platform endpoints, enabling end-to-end automated customer actions - Architected a concurrent multi-agent orchestration system using LangGraph and LangChain, automating complex agentic AI workflows and reducing end-to-end latency by 60% through optimized prompt routing - Engineered generative code solutions (API-aware templates and SDK patterns) to accelerate developer integration, ensuring high-ownership and rapid shipping of production-quality AI features - Implemented a robust observability and evaluation layer using LangSmith and Langfuse for tracing, structured logs, and evaluation hooks to maintain reliability in fast-paced startup environments - Optimized agent concurrency and KV-cache efficiency to maximize throughput, achieving a 50% reduction in LLM API costs while maintaining high execution performance	

AI Researcher Pace Artificial Intelligence Lab	New York City, NY Sep 2025 - Dec 2025
- Researched and deployed AI-native workflows using PyTorch, StyleGAN, and Stable Diffusion, leveraging CUDA and LORA for fine-tuning models on HPC clusters to generate industry-standard designs - Built and served LLM inference models using vLLM and Hugging Face, ensuring secure, low-latency delivery of production-ready AI features within a high-velocity research environment - Spearheaded workshops on Agentic AI (LangChain, LangGraph) and GPU programming for 50+ students, fostering a culture of technical excellence and knowledge sharing	

Software Engineer Qualitest	Noida, India Feb 2021 - Dec 2023
- Developed a full-stack support chatbot using Python, React, and RAG (Retrieval-Augmented Generation), reducing ticket resolution by 20% through automated AI-native products and workflows - Designed and developed a cloud-agnostic storage service (S3, GCS, MinIO) with Python and Redis, managing secure, signed-URL operations for production-scale traffic and data persistence - Contributed to an Identity & Access Management platform using React, Spring Boot, and AWS, supporting thousands of users with a focus on end-to-end reliability and execution mindset - Developed backend workflows using AWS Step Functions and Lambdas to handle long-running data import jobs, ensuring robust execution and preventing system timeouts	

PROJECTS

Dark Factory - Orchestrator for AI coding CLIs <i>Crossterm, Ratatui, Y-crdt</i>	github.com/ozzyozbourne/dark-factory
Built an event-driven orchestrator for AI coding agents using Crossterm and Ratatui, enabling secure, real-time command execution and distributed agentic loops	

- Implemented **Y-crdt** for signaling across distributed machines, allowing for full ownership of CLI IO and eliminating polling overhead in **AI-native** sessions