

Osaid Khan

[+1 \(945\) 304-5781](tel:+19453045781) | khanosaid726@gmail.com | www.linkedin.com/in/-osaid-khan/ | github.com/ozzyozbourne

Dear Hiring Team,

It is with great pleasure that I submit my candidacy for the Data Products Engineer role with Persona. I bring over three years of professional experience building customer-facing data products, backend services, APIs, and AI-powered agentic analytics tools, along with a strong academic foundation as a Master of Science in Computer Science graduate from Pace University, NY.

Through my experience across two engineering roles, I have built and delivered production-grade systems that directly reflect Persona's need for data-intensive products over large, fast-moving datasets. As a Software Engineer at Sperse, I built customer-facing data products in Python, LangGraph, LangChain, and FastAPI that let customers and internal teams ask questions of operational data in natural language, routing requests across 40+ specialized sub-agents with tenant-aware access control and PII handling for 10,000+ users. I also partnered with product and design to ship React Flow and TypeScript surfaces where users query, explore, export, and act on their data through agentic swarms, and created a secure FastMCP server that exposed 160+ backend actions through backend services, APIs, and a lightweight query layer for real-time, high-volume workflows. Alongside this work, I built OpenAI Codex code-review automation for 25-30 PRs/week and LlamaIndex-powered RAG agents with OpenAI embeddings and Pgvector as semantic layers over product documentation. Prior to Sperse, as a Software Engineer at Qualitest, I designed a cloud-agnostic Storage Service in Java and Spring Boot with PostgreSQL and Redis for uploads, downloads, deletions, and custom and bulk data exports across Google Cloud (GCP), S3, and MinIO via expiring signed URLs. I also optimized in-dashboard SQL workloads for analytical data stores using PostgreSQL window functions, materialized views, and indexes, and orchestrated jobs with AWS Step Functions and Lambda across 100K+ weekly executions. Together, these roles have prepared me to build Persona's customer-facing query, exploration, export, and analytics experiences with the performance, correctness, and security discipline this role requires.

Persona's mission to humanize online identity, build a safer and more accessible internet, and act as the identity layer that powers it strongly matches the way I approach engineering: make complex systems trustworthy, usable, and secure for the people depending on them. My Dark Factory project reflects that same bias toward disciplined agentic systems; I built distributed agentic tooling that turns large coding requests into dependency-aware task DAGs, then modeled execution state so agents and operators can query progress, inspect failures, and resume work safely. Persona's people-first and collaborative culture also resonates with my experience partnering with product and design to translate customer needs into reliable surfaces. I would be excited to bring that combination of backend depth, AI/LLM product experience, and data-platform rigor to Persona's Data Products team. I would love to discuss my interest in this position with your team and can be reached at +1 (945) 304-5781 or khanosaid726@gmail.com. Thank you for your time and consideration, I look forward to connecting!

Sincerely,
Osaid Khan