

Osaid Khan

+1 (945) 304-5781 | khanosaid726@gmail.com | www.linkedin.com/in/-osaid-khan/ | github.com/ozzyozbourne

EDUCATION

Pace University <i>Master of Science (MS) in Computer Science (GPA - 4.0/4.0)</i>	New York City, NY <i>Jan 2024 - Dec 2025</i>
Dr. A.P.J. Abdul Kalam Technical University <i>Bachelor of Technology (B.Tech) in Computer Science</i>	Noida, India <i>Jul 2016 - Sep 2020</i>

TECHNICAL SKILLS

Languages: Python, Java, Rust, JavaScript/TypeScript
Others: Kafka, Git, CI/CD, ROS, Node.js, Google Cloud
Databases: PostgreSQL, Redis, MongoDB
AWS Services: EC2, RDS, S3, Lambda, CloudWatch, Step Functions, CloudFormation, CloudFront, API Gateway
Frameworks: FastAPI, FastMCP, PyTorch, React, LangGraph, LangChain, LangSmith, Sprint Boot
AI: Claude code, Codex, Pi coding-agent

PROFESSIONAL EXPERIENCE

Software Engineer Sperse	Phoenix, AZ <i>Feb 2026 - present</i>
- Architected a concurrent multi-agent orchestration system using LangChain and LangGraph, enabling scalable, reliable, and performant tool execution across complex backend workflows - Developed an MCP interaction layer that translates natural-language intents into safe, validated tool calls across 1000+ platform endpoints, supporting task planning and automated actions - Built generative code solutions (API-aware templates and SDK patterns) to accelerate service design and integration with platform APIs, ensuring high code quality and system design standards - Enabled a robust observability layer using Langfuse and LangSmith for tracing, structured logs, and evaluation hooks to monitor system performance, reliability, and maintainability - Optimized prompt routing and agent concurrency to maximize KV-cache hit rate, reducing end-to-end latency by 60% and improving distributed system efficiency	

AI Researcher Pace Artificial Intelligence Lab	New York City, NY <i>Sep 2025 - Dec 2025</i>
- Conducted research on StyleGAN and Stable Diffusion to generate industry-standard designs using CUDA, PyTorch and LORA on an HPC cluster, managing complex concurrency and resource allocation - Utilized vLLM for model inference with Hugging Face deployment, ensuring secure and efficient model serving within a distributed system environment - Led technical workshops on GPU programming with CUDA, Mojo and building AI Agents, contributing to the team's engineering standards and knowledge-sharing	

Software Engineer Qualitest	Noida, India <i>Feb 2021 - Dec 2023</i>
- Designed and developed a cloud-agnostic Storage Service using Python, PostgreSQL, and Redis that handles real-world operational file operations across Google Cloud Storage, S3, and MinIO - Developed event-driven backend workflows using AWS Step Functions and Lambda to handle long-running data jobs, ensuring system reliability and preventing timeouts in distributed architectures - Built an internal support chatbot using LangChain and RAG with PostgreSQL vector store, enabling natural language queries over documentation and improving system maintainability - Contributed to an Identity & Access Management platform using React, Spring Boot, and AWS, supporting thousands of users and ensuring scalable service design	

PROJECTS

Dark Factory - Orchestrator for AI coding CLIs <i>Crossterm, Ratatui, Y-crdt</i>	github.com/ozzyozbourne/dark-factory
Enables full ownership of CLI IO via PTY, providing a secure interface with real-world shell environments and event-driven agentic loops	

- Implements a **distributed** orchestration layer using **Y-crdt** to coordinate task execution across multiple machines, ensuring **reliability** and **concurrency**