

Osaid Khan

+1 (945) 304-5781 | khanosaid726@gmail.com | www.linkedin.com/in/-osaid-khan/ | github.com/ozzyozbourne

EDUCATION

Pace University <i>Master of Science (MS) in Computer Science (GPA - 4.0/4.0)</i>	New York City, NY <i>Jan 2024 - Dec 2025</i>
Dr. A.P.J. Abdul Kalam Technical University <i>Bachelor of Technology (B.Tech) in Computer Science</i>	Noida, India <i>Jul 2016 - Sep 2020</i>

TECHNICAL SKILLS

Languages: Java, Python, JavaScript/TypeScript, Rust
AWS Services: EC2, RDS, S3, Lambda, CloudWatch, Step Functions, CloudFormation, CloudFront, API Gateway
Databases: PostgreSQL, MongoDB, Redis
Frameworks: Spring Boot, React, FastAPI, FastMCP, LangGraph, LangChain, LangSmith, PyTorch
Others: CI/CD, Git, Node.js, Google Cloud, Kafka, Technical Documentation, Unit Testing, Integration Testing
AI: Claude code, Codex, Pi coding-agent

PROFESSIONAL EXPERIENCE

Software Engineer Sperse	Phoenix, AZ <i>Feb 2026 - present</i>
------------------------------------	--

- Developed and implemented an MCP AI interaction layer that translates natural-language intents into **safe, validated tool calls** across **1000+ platform endpoints**, ensuring **robust software solutions** and automated customer actions
- Built **generative code solutions** and **API-aware templates** for developers to accelerate integration with platform **APIs**, improving **software development efficiency** and implementation correctness through **technical documentation**
- Enabled a robust **observability layer** using **Langfuse** and **LangSmith** for **tracing** and **structured logs** to **troubleshoot and resolve software defects**, monitoring behavior and tracking quality over time
- Architected a **concurrent multi-agent** orchestration system using **LangChain**, enabling scalable tool execution and reliable coordination across complex workflows with a focus on **security and compliance**
- Optimized prompt routing and agent concurrency to maximize **KV-cache hit rate**, reducing end-to-end latency by **60%** and ensuring **efficient software performance** while lowering operational costs

AI Researcher Pace Artificial Intelligence Lab	New York City, NY <i>Sep 2025 - Dec 2025</i>
--	---

- Conducted research on generative models using **PyTorch** on an **HPC cluster**, managing the full lifecycle from model training to inference optimization with **vLLM**, ensuring **performance and security**
- Conducted workshops on GPU programming and **AI Agents** for **+50** students, demonstrating strong **technical writing** and **communication skills** to articulate complex technical concepts clearly to diverse groups
- Applied **standard techniques and procedures** to manage model training runs, utilizing **LORA** for efficient fine-tuning while maintaining high-quality outputs for industry-standard applications

Software Engineer Qualitest	Noida, India <i>Feb 2021 - Dec 2023</i>
---------------------------------------	--

- Designed and developed a cloud-agnostic Storage Service using **Python**, **PostgreSQL**, and **Redis** that handles file operations in **AWS S3** and **Google Cloud Storage**, incorporating **unit testing and integration testing** for reliability
- Contributed to an Identity & Access Management platform using **Java**, **Spring Boot**, and **AWS**, supporting thousands of users while ensuring **software security and compliance** with industry standards
- Developed backend workflows using **AWS Step Functions** to manage long-running data jobs, performing **debugging and troubleshooting** to resolve complex issues and prevent system timeouts
- Collaborated with **cross-functional teams** to define and ship new features, providing **constructive feedback** during peer code reviews and maintaining comprehensive **technical documentation** for all program changes

PROJECTS

Dark Factory - Orchestrator for AI coding CLIs <i>Crossterm, Ratatui, Y-crdt</i>	github.com/ozzyozbourne/dark-factory
• Enables full ownership of CLI IO via PTY , implementing security measures to log every command before execution by the bash shell	

- Provides **event-driven** agentic loops for real-time updates, utilizing **analysis of data usage** and systems flow to eliminate polling overhead and improve performance