

Osaid Khan

+1 (945) 304-5781 | khanosaid726@gmail.com | www.linkedin.com/in/-osaid-khan/ | github.com/ozzyozbourne

EDUCATION

Pace University <i>Master of Science (MS) in Computer Science (GPA - 4.0/4.0)</i> University <i>Bachelor of Technology (B.Tech) in Computer Science</i>	<i>Jan 2024 - Dec 2025</i> Dr. A.P.J. Abdul Kalam Technical University <i>Jul 2016 - Sep 2020</i>	New York City, NY Noida, India
--	--	-----------------------------------

TECHNICAL SKILLS

Languages: Java, JavaScript/TypeScript, Rust, Python
Frameworks: FastAPI, FastMCP, LangGraph, LangChain, LangSmith, PyTorch, React, Sprint Boot
Databases: PostgreSQL, MongoDB, Redis
AWS Services: EC2, RDS, S3, Lambda, CloudWatch, Step Functions, CloudFormation, CloudFront, API Gateway
AI: Claude code, Codex, Pi coding-agent
Others: Google Cloud, Git, CI/CD, Kafka, Node.js

PROFESSIONAL EXPERIENCE

Software Engineer Spense	<i>Feb 2026 - present</i>	Phoenix, AZ
- Developed and delivered an MCP AI interaction layer that translates natural-language intents into safe, validated tool calls across 1000+ platform endpoints, enabling end-to-end automated customer actions - Built generative code solutions for developers (API-aware templates, SDK patterns, and code-generation workflows) to accelerate integration with platform APIs and improve the correctness of implementations - Enabled a robust observability layer using Langfuse and LangSmith for tracing, structured logs, and evaluation hooks to monitor behavior, debug failures, and track quality over time - Architected a concurrent multi-agent orchestration system using LangChain and LangGraph, enabling scalable tool execution and reliable coordination across complex customer workflows - Optimized prompt routing and agent concurrency to maximize KV-cache hit rate, reducing end-to-end latency by 60% and lowering LLM API costs by 50%		

AI Researcher Pace Artificial Intelligence Lab	<i>Sep 2025 - Dec 2025</i>	New York City, NY
- Conducted research on StyleGAN and Stable Diffusion to generate industry-standard fashion designs using CUDA, PyTorch and LORA (Low-rank Adaptation) on Pace University's HPC cluster for model training and fine-tuning - Utilized vLLM for model inference with Hugging Face deployment, ensuring secure and efficient model serving - Conducted workshops on GPU programming with CUDA, Mojo and building AI Agents with LangChain, LangGraph, and LangSmith for +50 students		

Software Engineer Qualitest	<i>Feb 2021 - Dec 2023</i>	Noida, India
- Built an internal support chatbot using LangChain and RAG (Retrieval-Augmented Generation) with PostgreSQL vector store, reducing average ticket resolution time by 20%, enabling natural language queries over documentation - Designed and developed a cloud-agnostic Storage Service using Python, PostgreSQL, and Redis that seamlessly handles file uploads, downloads, and deletions in Google Cloud Storage, S3, and MinIO through the utilization of expiring signed URLs - Contributed to the development of an Identity & Access Management platform using React, Spring Boot, and AWS, supporting thousands of users across 30+ applications - Developed backend workflows using AWS Step Functions to handle long-running data import jobs and prevent AWS Lambda timeouts		

PROJECTS

Dark Factory - Distributed orchestrator for coding agents <i>Elixir, TypeScript</i> github.com/ozzyozbourne/dark-factory
- Turns a single user goal into a dependency-aware task plan that is automatically dispatched to a fleet of AI coding agents running across multiple machines, with built-in retries and timeout handling - Built a fault-tolerant Elixir coordinator that matches tasks to agents by capability and live workload over a real-time WebSocket protocol, with a CRDT-backed distributed registry keeping the cluster alive when individual nodes fail